

# AGEC 652 - Lecture 5.1

## Unconstrained optimization

Part B: Line search and trust region methods

**Diego S. Cardoso**

**Spring 2023**

# Course roadmap

1. Intro to Scientific Computing
2. Numerical operations and representations
3. Systems of equations
4. Function approximation (Skipped)
5. **Optimization**
  - 5.1 **Unconstrained optimization**
    - A) Theory and derivative-free methods
    - B) **Line-search and trust region methods** ← *You are here*
  - 5.2 Constrained optimization
6. Structural estimation

\*These slides are based on Miranda & Fackler (2002), Nocedal & Wright (2006), Judd (1998), and course materials by Ivan Rudik and Florian Oswald.

# Solution strategies: line search vs. trust region

When we move from  $x^{(k)}$  to the next iteration,  $x^{(k+1)}$ , we have to decide

- Which direction from  $x^{(k)}$
- How far to go from  $x^{(k)}$

There are two fundamental solution strategies that differ in the order of those decisions

- **Line search** methods first choose a *direction* and then select the optimal *step size*
- **Trust region** methods first choose a *step size* and then select the optimal *direction*

# Line search algorithms

# Line search algorithms

General idea:

1. Start at some current iterate  $x_k$
2. Select a direction to move in  $p_k$
3. Figure out how far along  $p_k$  to move

# Line search algorithms

How do we figure out how far to move?

"Approximately" solve this problem to figure out the **step length**  $\alpha$

$$\min_{\alpha > 0} f(x_k + \alpha p_k)$$

We are finding the distance to move (  $\alpha$  ) along direction  $p_k$  that minimizes our objective  $f$

Typically, algorithms do not perform the full minimization problem since it is costly

- We only try a limited number of step lengths and stop when an approximation criterion is met (ex: Armijo, Wolfe, or Goldstein conditions)

# Line search: step length selection

Typical line search algorithms select the step length in two stages

1. Bracketing: pick an interval with desirable step lengths
2. Bisection or interpolation: find a "good" step length in this interval



# Line search: step length selection

A widely-used method is the **Backtracking** procedure

1. Choose  $\bar{\alpha} > 0, \rho \in (0, 1), c \in (0, 1)$
  2. Set  $\alpha \leftarrow \bar{\alpha}$
  3. Repeat until  $f(x_k + \alpha p_k) \leq f(x_k) + c\alpha \nabla f_k^T p_k$ 
    - $\alpha \leftarrow \rho\alpha$
  4. Terminate with  $\alpha_k = \alpha$
- Step 3 checks the *Armijo condition*, which checks for a *sufficient decrease* for convergence

\*Several other step length methods exist. See Nocedal & Wright Ch.3 and Miranda & Fackler Ch 4.4 for more examples.

# Line search: direction choice

**We still haven't answered, what direction  $p_k$  do we decide to move in?**


What's an obvious choice for  $p_k$ ?

The direction that yields the *steepest descent*

—  $-\nabla f_k$  is the direction that makes  $f$  decrease most rapidly

- $k$  indicates we are evaluating  $f$  at iteration  $k$

# Steepest descent method



S

# Steepest descent method

We can verify this is the direction of steepest descent by referring to Taylor's theorem

For any direction  $p$  and step length  $\alpha$ , we have that

$$f(x_k + \alpha p) = f(x_k) + \alpha p^T \nabla f_k + \frac{1}{2!} \alpha^2 p^T \nabla^2 f(x_k + tp) p$$

The rate of change in  $f$  along  $p$  at  $x_k$  ( $\alpha = 0$ ) is  $p^T \nabla f_k$

# Steepest descent method

The the unit vector of quickest descent solves

$$\min_p p^T \nabla f_k \quad \text{subject to: } \|p\| = 1$$

Re-express the objective as

$$\min_{\theta, \|p\|} \|p\| \|\nabla f_k\| \cos \theta$$

where  $\theta$  is the angle between  $p$  and  $\nabla f_k$

The minimum is attained when  $\cos \theta = -1$  and  $p = -\frac{\nabla f_k}{\|\nabla f_k\|}$ , so the direction of steepest descent is simply  $-\nabla f_k$

# Steepest descent method

The **steepest descent method** searches along this direction at every iteration  $k$


- It may select the step length  $\alpha_k$  in several different ways
- A benefit of the algorithm is that we only require the gradient of the function, and no Hessian
- However it can be very slow

# Line search: alternative directions

We can always use search directions other than the steepest descent

Any descent direction (i.e. one with angle strictly less than  $90^\circ$  of  $-\nabla f_k$ ) is *guaranteed* to produce a decrease in  $f$  as long as the step size is sufficiently small

**But is  $-\nabla f_k$  always the best search direction?**



# Newton-Raphson method

The most important search direction is not steepest descent but **Newton's direction**

This direction gives rise to the Newton-Raphson Method

- This method is basically just using Newton's method to find the root of the gradient of the objective function



# Newton-Raphson method

Newton's direction comes out of the second order Taylor series approximation to  $f(x_k + p)$

$$f(x_k + p) \approx f_k + p^T \nabla f_k + \frac{1}{2!} p^T \nabla^2 f_k p$$

We find the Newton direction by selecting the vector  $p$  that minimizes  $f(x_k + p)$

This ends up being

$$p_k^N = -[\nabla^2 f_k]^{-1} \nabla f_k$$

# Newton-Raphson method

The algorithm is pretty much the same as in Newton's rootfinding method

1. Start with an initial guess  $x_0$
  2. Repeat until convergence
    - $x_{k+1} \leftarrow x_k - \alpha_k [\nabla^2 f_k]^{-1} \nabla f_k$ 
      - where  $\alpha_k$  comes from a step length selection algorithm
  3. Terminate with  $x^* = x_k$
- Most packages just use  $\alpha = 1$  (i.e., Newton's method step). But you can usually change this parameter if you have convergence issues

# Newton-Raphson method

This approximation to the function we are trying to solve has error of  $O(\|p\|^3)$ , so if  $p$  is small, the quadratic approximation is very accurate

## Drawbacks:

- The Newton direction is only guaranteed to decrease the objective function if  $\nabla^2 f_k$  is positive definite
- It requires explicit computation of the Hessian,  $\nabla^2 f(x)$ 
  - But quasi-Newton solvers also exist

# Quasi-Newton methods

Just like in rootfinding, there are several methods to avoid computing derivatives (Hessians, in this case)

Instead of the true Hessian  $\nabla^2 f(x)$ , these methods use an approximation  $B_k$  (to the inverse of the Hessian). Hence, they set direction

$$d_k = -B_k \nabla f_k$$

The optimization method analogous to Broyden's that also uses the *secant condition* is the **BFGS method**

- Named after its inventors, Broyden, Fletcher, Goldfarb, Shanno

# Linear search methods in Julia

Once again, we will use `Optim.jl`. We'll see an example with an easy function, solving it using Steepest Descent, Newton-Raphson, and BFGS


$$f(x_1, x_2) = ax_1^2 + bx_2^2 + cx_1 + dx_2 + ex_1x_2$$

We will use parameters  $a = 1, b = 4, c = -2, d = -1, e = -3$

```
using Optim, Plots, LinearAlgebra;
a = 1; b = 4; c = -2; d = -1; e = -3;
f(x) = a*x[1]^2 + b*x[2]^2 + c*x[1] + d*x[2] + e*x[1]*x[2];
```

# Linear search methods in Julia

Let's take a look at our function with a contour plot



# Linear search methods in Julia

Since we will use Newton-Raphson, we should define the gradient and the Hessian of our function

```
# Gradient
function g!(G, x)
    G[1] = 2a*x[1] + c + e*x[2]
    G[2] = 2b*x[2] + d + e*x[1]
end;

# Hessian
function h!(H, x)
    H[1,1] = 2a
    H[1,2] = e
    H[2,1] = e
    H[2,2] = 2b
end;
```

# Linear search methods in Julia

Let's check if the Hessian satisfies it being positive semidefinite. One way is to check whether all eigenvalues are positive. In this case,  $H$  is constant, so it's easy to check

```
H = zeros(2,2);  
h!(H, [0 0]);  
LinearAlgebra.eigen(H).values
```

```
## 2-element Vector{Float64}:  
##  0.7573593128807148  
##  9.242640687119286
```



# Linear search methods in Julia

Since the gradient is linear, it is also easy to calculate the minimizer analytically. The FOC is just a linear equation

```
analytic_x_star = [2a e; e 2b]\[-c ;-d]
```

```
## 2-element Vector{Float64}:  
##  2.714285714285715  
##  1.142857142857143
```

# Linear search methods in Julia

Let's solve it with the Steepest (or Gradient) descent method

```
# Initial guess
x0 = [0.2, 1.6];
res_GD = Optim.optimize(f, g!, x0, GradientDescent(), Optim.Options(x_abstol=1e-3))
```

```
## * Status: success
##
## * Candidate solution
##   Final objective value:      -3.285714e+00
##
## * Found with
##   Algorithm:      Gradient Descent
##
## * Convergence measures
##    $|x - x'|$  = 4.55e-04  $\leq$  1.0e-03
##    $|x - x'|/|x'|$  = 1.68e-04  $\leq$  0.0e+00
##    $|f(x) - f(x')|$  = 1.25e-06  $\leq$  0.0e+00
##    $|f(x) - f(x')|/|f(x')|$  = 3.80e-07  $\leq$  0.0e+00
##    $|g(x)|$  = 4.83e-04  $\leq$  1.0e-08
##
## * Work counters
```

# Linear search methods in Julia

Let's solve it with the Steepest (or Gradient) descent method

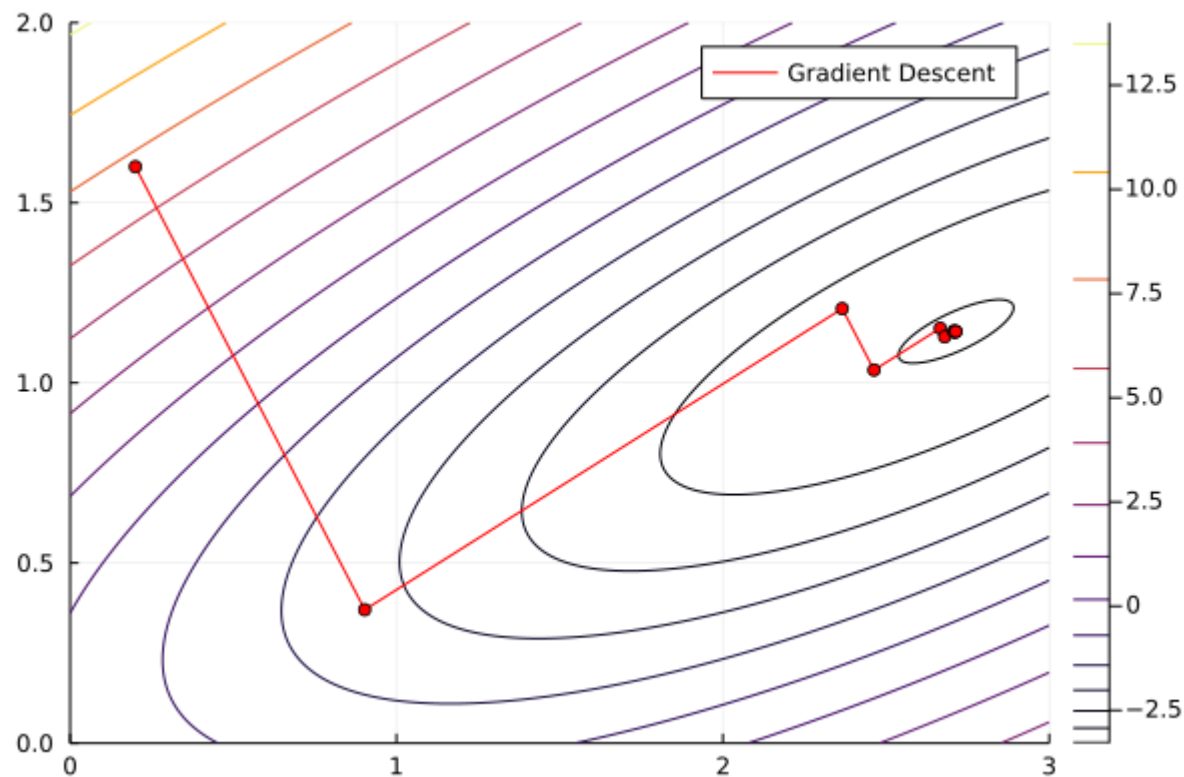
```
res_GD.minimizer
```

```
## 2-element Vector{Float64}:  
##  2.7136160576693604  
##  1.1425715540060508
```

```
res_GD.minimum
```

```
## -3.285714084769726
```

# Linear search methods in Julia



# Linear search methods in Julia

We haven't really specified a line search method yet

In most cases, `Optim.jl` will use by default the *Hager-Zhang method*

- This is based on Wolfe conditions

But we can specify other approaches. We need the `LineSearches` package to do that:

```
using LineSearches;
```

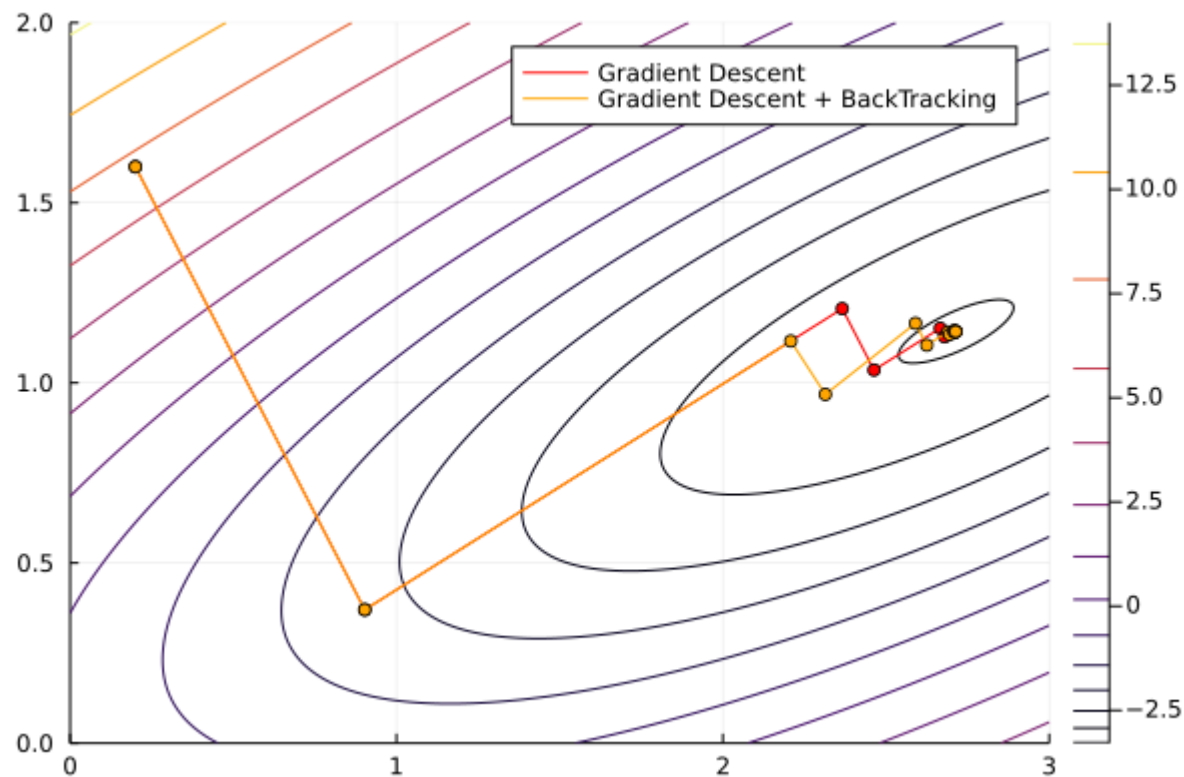
Let's re-run the `GradientDescent` method using  $\bar{\alpha} = 1$  and the backtracking method

# Linear search methods in Julia

```
Optim.optimize(f, g!, x0,  
               GradientDescent(alphaguess = LineSearches.InitialStatic(alpha=1.0),  
                               linesearch = BackTracking()),  
               Optim.Options(x_abstol=1e-3))
```

```
## * Status: success  
##  
## * Candidate solution  
##   Final objective value:      -3.285714e+00  
##  
## * Found with  
##   Algorithm:      Gradient Descent  
##  
## * Convergence measures  
##    $|x - x'|$  =  $3.61e-04 \leq 1.0e-03$   
##    $|x - x'|/|x'|$  =  $1.33e-04 \leq 0.0e+00$   
##    $|f(x) - f(x')|$  =  $8.65e-07 \leq 0.0e+00$   
##    $|f(x) - f(x')|/|f(x')|$  =  $2.63e-07 \leq 0.0e+00$   
##    $|g(x)|$  =  $6.73e-04 \leq 1.0e-08$   
##  
## * Work counters  
##   Seconds run:  0 (vs limit Inf)  
##   Iterations:  11
```

# Linear search methods in Julia



# Linear search methods in Julia

Next, let's use the Newton-Raphson method with default (omitted) line search parameters


- If you omit `g!` and `h!`, `Optim` will approximate them numerically for you. You can also specify options to use auto differentiation

```
Optim.optimize(f, g!, h!, x0, Newton())
```

```
## * Status: success
##
## * Candidate solution
##   Final objective value:      -3.285714e+00
##
## * Found with
##   Algorithm:      Newton's Method
##
## * Convergence measures
##   |x - x'|          = 2.51e+00 ≤/0.0e+00
##   |x - x'|/|x'|     = 9.26e-01 ≤/0.0e+00
##   |f(x) - f(x')|    = 1.06e+01 ≤/0.0e+00
```



# Linear search methods in Julia



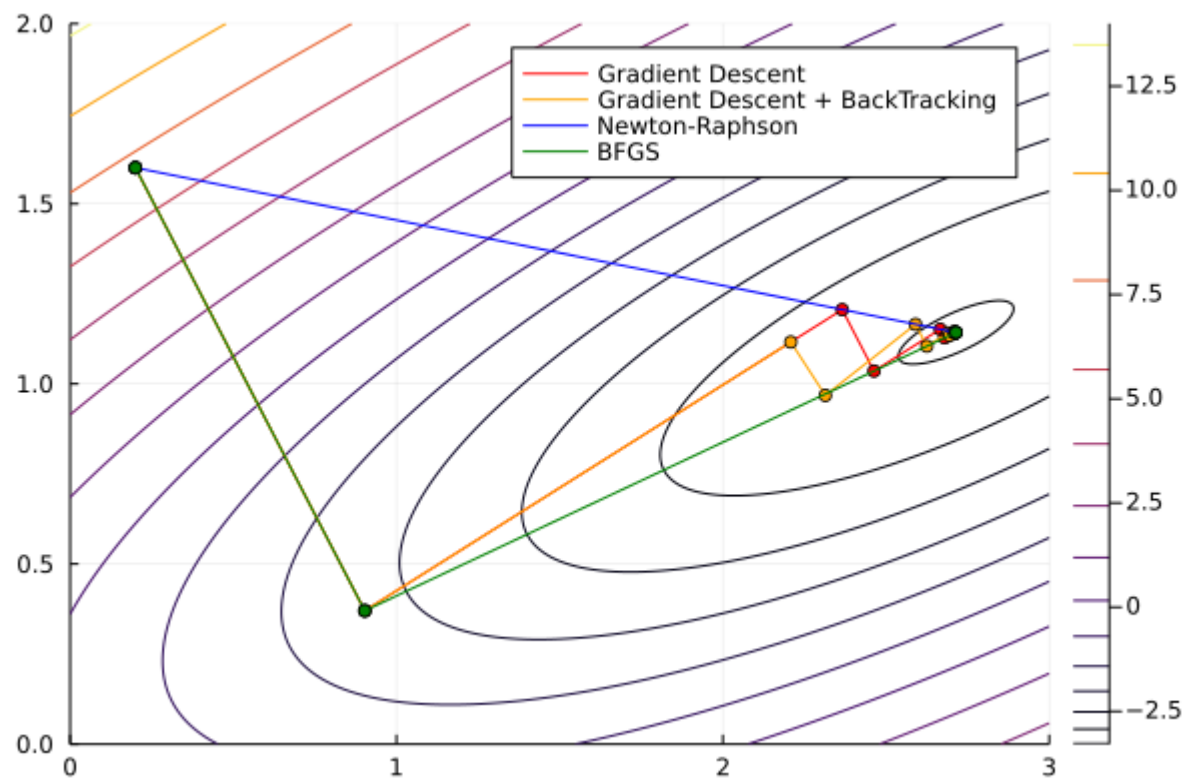
# Linear search methods in Julia

Lastly, let's use the BFGS method

```
Optim.optimize(f, x0, BFGS())
```

```
## * Status: success
##
## * Candidate solution
##   Final objective value:      -3.285714e+00
##
## * Found with
##   Algorithm:      BFGS
##
## * Convergence measures
##    $|x - x'|$  = 1.81e+00  $\leq$  0.0e+00
##    $|x - x'|/|x'|$  = 6.67e-01  $\leq$  0.0e+00
##    $|f(x) - f(x')|$  = 1.47e+00  $\leq$  0.0e+00
##    $|f(x) - f(x')|/|f(x')|$  = 4.48e-01  $\leq$  0.0e+00
##    $|g(x)|$  = 1.28e-10  $\leq$  1.0e-08
##
## * Work counters
##   Seconds run:    0 (vs limit Inf)
##   Iterations:    2
```


# Linear search methods in Julia



# Trust regions algorithms

# Trust region methods

Trust region methods construct an approximating model,  $m_k$  whose behavior near the current iterate  $x_k$  is close to that of the actual function  $f$



We then search for a minimizer of  $m_k$

# Trust region methods

**Issue:**  $m_k$  may not represent  $f$  well when far away from the current iterate  $x_k$

**Solution:** Restrict the search for a minimizer to be within some region of  $x_k$ , called a **trust region**

We are only going to cover the basic of trust region methods. For details, see Nocedal & Wright (2006), Chapter 4.

# Trust region methods

Trust region problems can be formulated as

$$\min_p m_k(x_k + p)$$

where  $x_k + p \in \Gamma$

- $\Gamma$  is a ball defined by  $\|p\|_2 \leq \Delta_k$
- $\Delta_k$  is called the trust region radius

$\Delta_k$  is adjusted every iteration based on how well  $m_k$  approximates  $f_k$  around current guess  $x_k$


# Trust region methods

Typically the approximating model  $m_k$  is a quadratic function (i.e. a second-order Taylor approximation)

$$m_k(x_k + p) = f_k + p^T \nabla f_k + \frac{1}{2!} p^T B_k p$$

where  $B_k$  is the Hessian or an approximation to the Hessian

Solving this problem usually involves finding the *Cauchy point*





# Trust region methods

From  $x_k$ , the Cauchy point can be found in the direction

$$p_k^C = -\tau_k \frac{\Delta_k}{\|\nabla f_k\|} \nabla f_k$$

So it's kind of a gradient descent ( $-\nabla f_k$ ), but with an adjusted step size within the trust region. The step size depends on the radius  $\Delta_k$  and parameter  $\tau_k$

$$\tau_k = \begin{cases} 1 & \text{if } \nabla f_k^T B_k \nabla f_k \leq 0 \\ \min(\|\nabla f_k\|^3 / \Delta_k \nabla f_k^T B_k \nabla f_k, 1) & \text{otherwise} \end{cases}$$




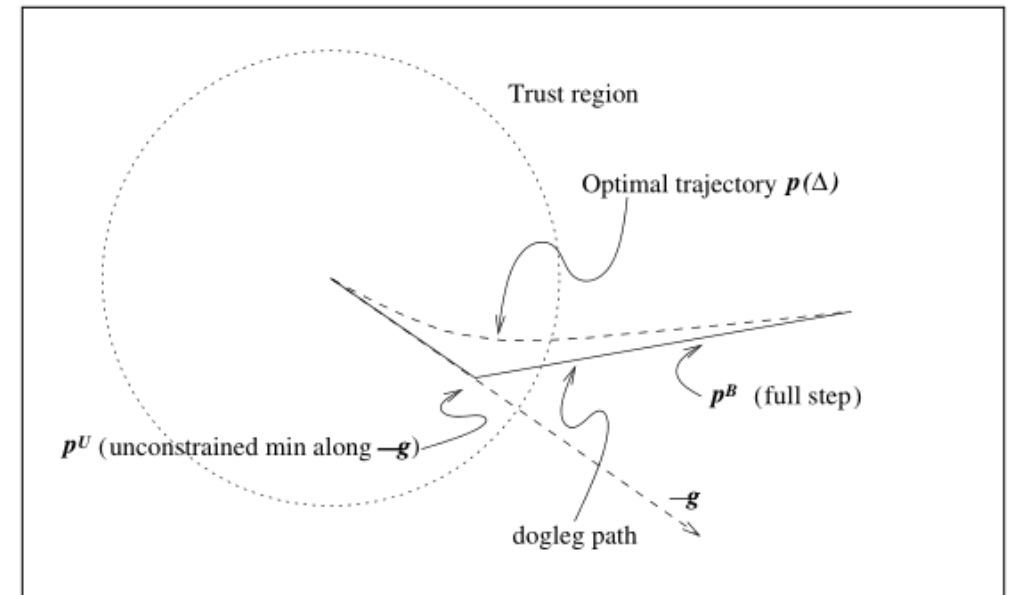
Figure 4.3 The Cauchy point.

# Trust region methods

If you ran `nlsolve` with default parameters, you may have noticed it uses **Trust region with dogleg**. What's the deal with the **dogleg**?

It's an improvement on the Cauchy point method

- It allows us to move in V-shaped trajectory instead of slowly adjusting with Cauchy directions along a curved path



**Figure 4.4** Exact trajectory and dogleg approximation.

# Trust region method illustration


Let's see an illustration of the trust region method with the Branin function (it has 3 global minimizers)\*

$$\min f(x_1, x_2) = (x_2 - 0.129x_1^2 + 1.6x_1 - 6)^2 + 6.07 \cos(x_1) + 10$$

- Source for this example:  
[https://optimization.mccormick.northwestern.edu/index.php/Trust-region\\_methods](https://optimization.mccormick.northwestern.edu/index.php/Trust-region_methods)


# Trust region method illustration

Contour plot




# Trust region method illustration

Iteration 1




# Trust region method illustration

Iteration 2




# Trust region method illustration

Iteration 3



# Trust region method illustration

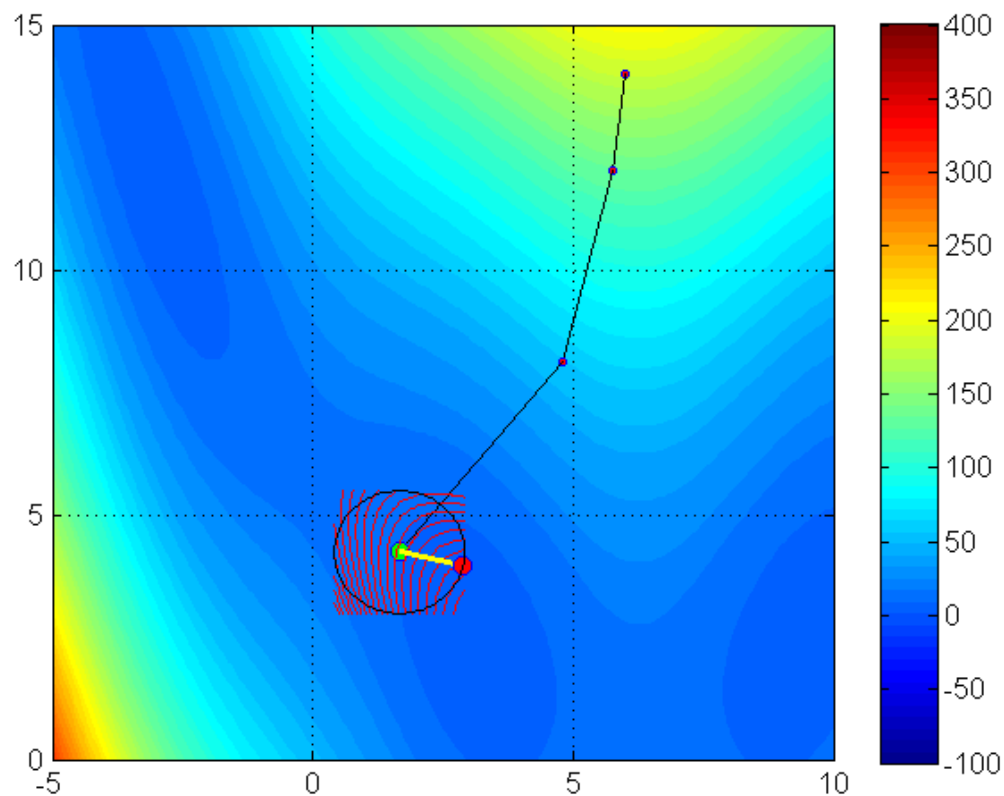
Iteration 4






# Trust region method illustration

Iteration 5




# Trust region method illustration

Iteration 6



# Trust region method illustration

Full trajectory



# Trust region methods in Julia

`Optim` includes a `NewtonTrustRegion()` solution method. To use it, you can call `optimize` just like you did with `Newton()`

# Your turn: Trust region methods in Julia

Use `Optim` and `NewtonTrustRegion*` to solve the same quadratic function from before

$$f(x_1, x_2) = ax_1^2 + bx_2^2 + cx_1 + dx_2 + ex_1x_2$$

With parameters  $a = 1, b = 4, c = -2, d = -1, e = -3$

```
a = 1; b = 4; c = -2; d = -1; e = -3;
f(x) = a*x[1]^2 + b*x[2]^2 + c*x[1] + d*x[2] + e*x[1]*x[2];
# Complete the arguments of the call
Optim.optimize(<?>, <?>, <?>, <?>, <?>)
```

\*Use default parameters. You can tweak solver parameters later by checking out the documentation  
[https://juliansolvers.github.io/Optim.jl/stable/#algo/newton\\_trust\\_region/](https://juliansolvers.github.io/Optim.jl/stable/#algo/newton_trust_region/)

# Trust region methods in Julia

```
## * Status: success
##
## * Candidate solution
##   Final objective value:      -3.285714e+00
##
## * Found with
##   Algorithm:      Newton's Method (Trust Region)
##
## * Convergence measures
##    $|x - x'|$  = 1.85e+00  $\leq$  0.0e+00
##    $|x - x'|/|x'|$  = 6.83e-01  $\leq$  0.0e+00
##    $|f(x) - f(x')|$  = 2.15e+00  $\leq$  0.0e+00
##    $|f(x) - f(x')|/|f(x')|$  = 6.54e-01  $\leq$  0.0e+00
##    $|g(x)|$  = 8.88e-16  $\leq$  1.0e-08
##
## * Work counters
##   Seconds run:    0   (vs limit Inf)
##   Iterations:    2
##   f(x) calls:    3
##    $\nabla f(x)$  calls: 3
##    $\nabla^2 f(x)$  calls: 3
```

# Trust region methods in Julia

